

Custom object training and detection with YOLOv3, Darknet and OpenCV



Vino Mahendran [Follow](#)

Nov 15, 2019 · 5 min read



Photo by Jessica Ruscello on Unsplash

YOLO is a state-of-the-art, real-time object detection system. It looks at the whole image at test time so its predictions are informed by global context in the image. It also makes predictions with a single network evaluation which makes it extremely fast when compared to R-CNN and Fast R-CNN.

This paper gives more details about how YOLO achieves the performance improvement.

Darknet

We will use Darknet, an open source neural network framework to train the detector. Download and build darknet

```
git clone https://github.com/pjreddie/darknet
cd darknet
make
```

Once that's successful, To test the build we can download pre trained YOLO weights and perform detection with the test image.

```
./darknet detector test cfg/coco.data cfg/yolov3.cfg
yolov3.weights data/dog.jpg
```


For training with custom objects, let us create the following

required files and directories

```
custom_data/custom.names  
custom_data/images  
custom_data/train.txt  
custom_data/test.txt
```

In this example we will train to detect Tesla cars for which we should collect the images and place it in `custom_data/images` directory.

Labels of our objects should be saved in `custom_data/custom.names` file, each line in the file corresponds to an object. In our case since we have only one object class, the file should contain the following



Annotation

After we collect the images containing our custom object, we will

need to annotate them. For YOLOv3, each image should have a corresponding text file with the same file name as that of the image in the same directory.

In our case text files should be saved in `custom_data/images` directory. For e.g. `image1.jpg` should have a text file `image1.txt`.

Each row in the text file corresponds to a single bounding box of the object and should have the following information

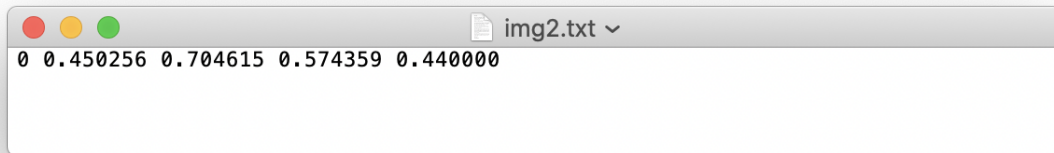
```
<object-class-id> <x-centre> <y-centre> <width> <height>
```

where


<object-class-id> an integer from 0 to (classes - 1) corresponding to the classes in the `custom_data/custom.names` file

height, width – Actual height and width of the image
x, y – centre coordinates of the bounding box
h, w – height and width of the bounding box

<x-centre> : x / width
<y-centre> : y / height
<width> : w / width
<height> : h / height



This can also be done with Labellmg, a graphical image annotation tool which creates `.txt` files for the images in YOLO format.




Train and Test sets

We can then randomly split the annotated images into train and

test sets in the ratio of 80:20

`custom_data/train.txt` Each row in the file should have the location of train dataset.

`custom_data/test.txt` Each row in the file should have the location of test dataset.



Pre-trained weights

To train our object detector we can use the existing pre trained weights that are already trained on huge data sets. From here we can download the pre trained weights to the root directory.

YOLO data file

Create a `detector.data` file in the `custom_data` directory which

should contain information regarding the train and test data sets

```
classes=1
train=custom_data/train.txt
valid=custom_data/test.txt
names=custom_data/custom.names
backup=backup/
```

backup is the location where newly trained weights would be saved.

Configurations

Based on the required performance we can select the YOLOv3 configuration file. For this example we will be using `yolov3.cfg`.

We can duplicate the file from `cfg/yolov3.cfg` to

```
custom_data/cfg/yolov3-custom.cfg
```

While training the images, weights of the neural networks are updated iteratively. We may use huge training sets which makes it resource consuming to update the weights for the entire training set in a single iteration. To use a small set of images to iteratively update the weights, `batch` param is set. By default it is set to 64.

The maximum number of iterations for which our network should be trained is set with the param `max_batches=4000`. Also update `steps=3200, 3600` which is 80%, 90% of `max_batches`.

We will need to update the `classes` and `filters` params of `[yolo]` and `[convolutional]` layers that are just before the `[yolo]` layers.

In this example since we have a single class (tesla) we will update the `classes` param in the `[yolo]` layers to 1 at line numbers:

610, 696, 783

Similarly we will need to update the `filters` param based on the classes count $\text{filters} = (\text{classes} + 5) * 3$. For a single class we should set `filters=18` at line numbers: 603, 689, 776

All the configuration changes are made to `custom_data/cfg/yolov3-custom.cfg`

Training

With all the required files and annotated images we can start our training

```
./darknet detector train custom_data/detector.data
custom_data/cfg/yolov3-custom.cfg darknet53.conv.74
```

```
→ darknet git:(master) * ./darknet detector train custom_data/detector.data custom_data/cfg/yolov3-custom.cfg darknet53.conv.74
yolov3-custom
layer    filters  size      input           output
0 conv   32  3 x 3 / 1  608 x 608 x 3  ->  608 x 608 x 32  0.639 BFLOPs
1 conv   64  3 x 3 / 2  608 x 608 x 32  ->  304 x 304 x 64  3.407 BFLOPs
2 conv   32  1 x 1 / 1  304 x 304 x 64  ->  304 x 304 x 32  0.379 BFLOPs
3 conv   64  3 x 3 / 1  304 x 304 x 32  ->  304 x 304 x 64  3.407 BFLOPs
```


4	res	1			304 x 304 x 64	->	304 x 304 x 64	
5	conv	128	3 x 3 / 2		304 x 304 x 64	->	152 x 152 x 128	3.407 BFLOPs
6	conv	64	1 x 1 / 1		152 x 152 x 128	->	152 x 152 x 64	0.379 BFLOPs
7	conv	128	3 x 3 / 1		152 x 152 x 64	->	152 x 152 x 128	3.407 BFLOPs
8	res	5			152 x 152 x 128	->	152 x 152 x 128	
9	conv	64	1 x 1 / 1		152 x 152 x 128	->	152 x 152 x 64	0.379 BFLOPs
10	conv	128	3 x 3 / 1		152 x 152 x 64	->	152 x 152 x 128	3.407 BFLOPs
11	res	8			152 x 152 x 128	->	152 x 152 x 128	
12	conv	256	3 x 3 / 2		152 x 152 x 128	->	76 x 76 x 256	3.407 BFLOPs
13	conv	128	1 x 1 / 1		76 x 76 x 256	->	76 x 76 x 128	0.379 BFLOPs
14	conv	256	3 x 3 / 1		76 x 76 x 128	->	76 x 76 x 256	3.407 BFLOPs
15	res	12			76 x 76 x 256	->	76 x 76 x 256	
16	conv	128	1 x 1 / 1		76 x 76 x 256	->	76 x 76 x 128	0.379 BFLOPs
17	conv	256	3 x 3 / 1		76 x 76 x 128	->	76 x 76 x 256	3.407 BFLOPs
18	res	15			76 x 76 x 256	->	76 x 76 x 256	
19	conv	128	1 x 1 / 1		76 x 76 x 256	->	76 x 76 x 128	0.379 BFLOPs
20	conv	256	3 x 3 / 1		76 x 76 x 128	->	76 x 76 x 256	3.407 BFLOPs
21	res	18			76 x 76 x 256	->	76 x 76 x 256	

We can continue training until the loss reaches a certain threshold. By default, weights for the custom detector is saved for every 100 iterations until 1000 iterations and then continues to save for every 10000 iterations. This behaviour can be modified by updating the condition at line 138 of `examples/detector.c` file.

Once the training is complete we can use the generated weights to perform detection.

Detection with OpenCV

We can perform detection with OpenCV DNN as it is a fast DNN implementation for CPU.

```

1  import argparse
2
3  import cv2
4  import numpy as np
5
6  parser = argparse.ArgumentParser(add_help=False)
7  parser.add_argument("--image", default='samples/image.jpg', help="image for
8  parser.add_argument("--config", default='cfg/yolov3.cfg', help="YOLO config
9  parser.add_argument("--weights", default='yolov3.weights', help="YOLO weight

```

```
10 parser.add_argument("--names", default='data/coco.names', help="class names")
11 args = parser.parse_args()
12
13 CONF_THRESH, NMS_THRESH = 0.5, 0.5
14
15 # Load the network
16 net = cv2.dnn.readNetFromDarknet(args.config, args.weights)
17 net.setPreferableBackend(cv2.dnn.DNN_BACKEND_OPENCV)
18 net.setPreferableTarget(cv2.dnn.DNN_TARGET_CPU)
19
20 # Get the output layer from YOLO
21 layers = net.getLayerNames()
22 output_layers = [layers[i[0] - 1] for i in net.getUnconnectedOutLayers()]
23
24 # Read and convert the image to blob and perform forward pass to get the bounding boxes
25 img = cv2.imread(args.image)
26 height, width = img.shape[:2]
27
28 blob = cv2.dnn.blobFromImage(img, 0.00392, (416, 416), swapRB=True, crop=False)
29 net.setInput(blob)
30 layer_outputs = net.forward(output_layers)
31
32 class_ids, confidences, b_boxes = [], [], []
33 for output in layer_outputs:
34     for detection in output:
35         scores = detection[5:]
36         class_id = np.argmax(scores)
37         confidence = scores[class_id]
38
39         if confidence > CONF_THRESH:
40             center_x, center_y, w, h = (detection[0:4] * np.array([width, height, width, height]).astype(int)).astype(int)
41
42             x = int(center_x - w / 2)
43             y = int(center_y - h / 2)
44
45             b_boxes.append([x, y, int(w), int(h)])
46             confidences.append(float(confidence))
47             class_ids.append(int(class_id))
48
```

```
49 # Perform non maximum suppression for the bounding boxes to filter overlapping
50 indices = cv2.dnn.NMSBoxes(b_boxes, confidences, CONF_THRESH, NMS_THRESH).flatten()
51
52 # Draw the filtered bounding boxes with their class to the image
53 with open(args.names, "r") as f:
54     classes = [line.strip() for line in f.readlines()]
55     colors = np.random.uniform(0, 255, size=(len(classes), 3))
56
57 for index in indices:
```

We can specify `--image`, `--config`, `--weights` and `--names` params as per our training to perform predictions for our custom object.

. . .

Francium Tech is a technology company laser focused on delivering top quality software of scale at extreme speeds. Numbers and Size of the data don't scare us. If you have any requirements or want a free health check of your systems or architecture, feel free to shoot an email to contact@francium.tech, we will get in touch with you!

Yolov3 Darknet Opencv Deep Learning Object Detection

Discover Medium

Welcome to a place where words matter. On Medium, smart voices and original ideas take center stage - with no ads in sight. Watch

Make Medium yours

Follow all the topics you care about, and we'll deliver the best stories for you to your homepage and inbox. Explore

Become a member

Get unlimited access to the best stories on Medium — and support writers while you're at it. Just \$5/month. Upgrade

[About](#)[Help](#)[Legal](#)